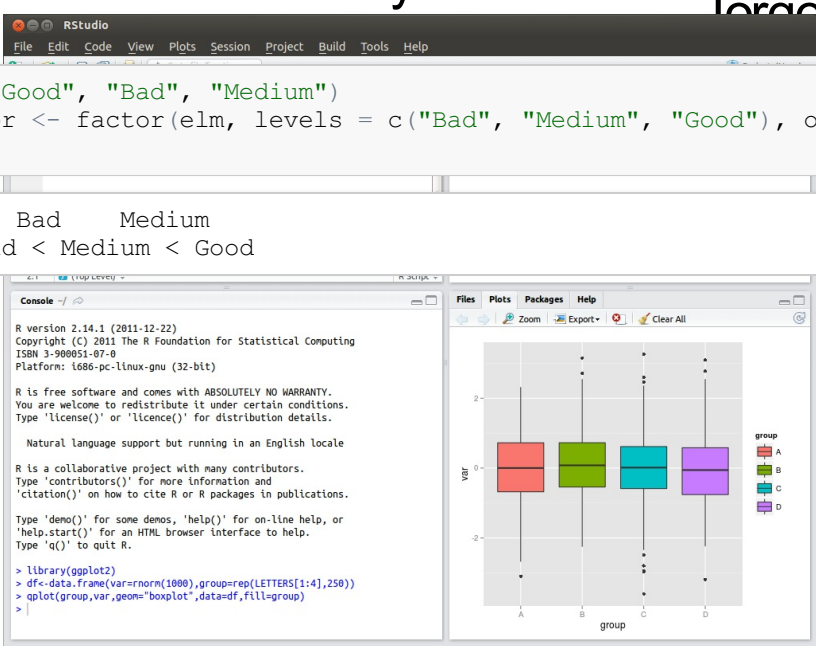


A factor variable is just a nice way of saying that a variable has unique values! And they can be ordered. This has consequences

```
elm <- c("Good", "Bad", "Medium")
(elm_factor <- factor(elm, levels = c("Bad", "Medium", "Good"), ordered = T))
```

[1] Good Bad Medium  
Levels: Bad < Medium < Good



```
table(elm)
```

elm	Bad	Good	Medium
	1	1	1

```
table(elm_factor)
```

elm_factor	Bad	Medium	Good
	1	1	1

How to install R?

Luckily, you guys have R and Rstudio installed, so you don't have to worry about this!

But if you want to install it at home, please follow [this guide](#)

That guide can help you install

- R
- Rstudio
- And `swirl`, a package in which you could do a bunch of exercises as homework!

What is R?

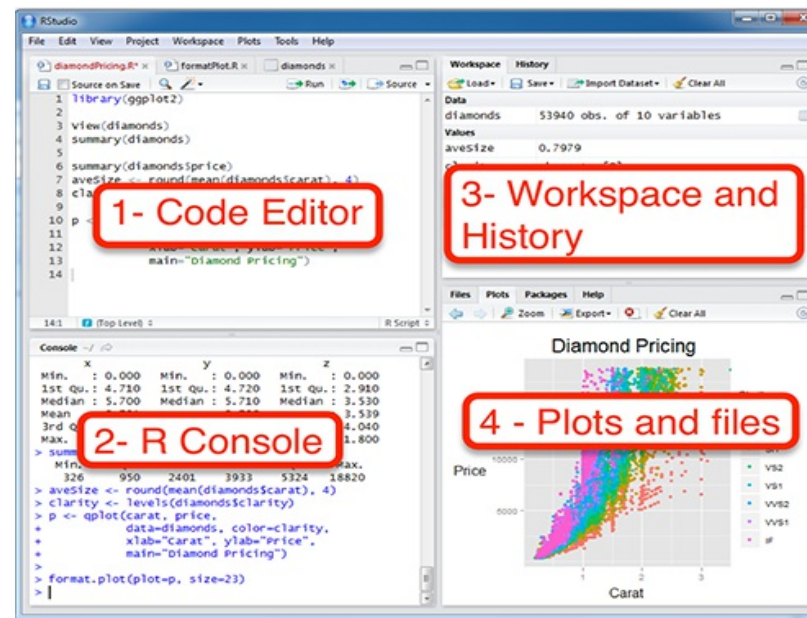
R is a programming language designed to do data analysis, usually interactive.

R is helpful for..

- Getting that darn excel/stata file into R (importing)
  - Turning that very ugly dataset into something to work with (data cleaning)
  - Automating your weekly reports (automating tasks)
  - Analyzing data (modeling)
  - Creating nicely formatted documents (communicating results)
  - Building your own commands to do specific things (functions)
- 
- Building very creative graphics
- 
- Among many things...

And so.. what is Rstudio?

And so.. what is Rstudio?



Let's get to it then!

R is an interactive language. That means that if you type a number,  
you will get a number.

```
#Input  
10
```

```
[1] 10
```

```
#Input  
5
```

```
[1] 5
```

## Introduction to R objects

R is also a calculator

Try typing these operations in R:

- `5 + 5`
- `10 - 5`
- `10 * 5`
- `20 / 10`
- `(10 * 20) - 5 / 2 + 2`
- `2 ^ 3`

Before we continue, what type of operations are these?

Answers in next slide!

## Introduction to R objects

- Addition
- Subtraction
- Multiplication
- Division
- A combination of all
- Exponentiation

Numbers in R are called `numeric`s.

For example:

```
is.numeric(10)  
is.numeric(10 + 20)  
is.numeric(10 / 2)
```



## Introduction to R objects

Having single numbers, like 10, is not very useful.

We want something similar to a column of a dataset, like age or income.

We can do that with `c()`, which stands for concatenate.

```
c(32, 34, 18, 22, 65)
```

```
[1] 32 34 18 22 65
```

Read this expression as: concatenate these numbers into a single object.

## Introduction to R objects

We can also give it a name, like age.

```
age <- c(32, 34, 18, 22, 65)
```

- Why didn't the result get printed?
- Where is this age object at?
- What is formally the age object?

## Introduction to R objects

We just created our first variable! The typical SAS/Excel/Stata column.

In R, these objects are called '**vectors**'.

Vectors can have several flavours:

- Numerics (we just saw one)
- Logicals
- Characters
- Factors

## Introduction to R objects

Suppose these ages belong to certain people. We can create a character vector with their names.

Following this guideline, create it yourself.

- Create a character vector with `c()`
- Include the names Paul, Maria, Andres, Roberto and Alicia inside
- `wrap` every name in quotes like this “Paul”, “Maria”, etc... This will make R understand that input as characters.

## Introduction to R objects

Answer:

```
c("Paul", "Maria", "Andres", "Robert", "Alicia")
```

```
[1] "Paul" "Maria" "Andres" "Robert" "Alicia"
```

We can also give it a name, like participants.

```
participants <- c("Paul", "Maria", "Andres", "Robert", "Alicia")
```

## Introduction to R objects

Character vectors are filled by strings, like “Paul” or “Maria”.

Can we do operations with strings?

```
"Paul" + "Maria"
```

```
Error in "Paul" + "Maria": argumento no-numérico para operador binario
```

Makes sense.. we can't add any letters.

Alright, we're set. Concatenate the numeric vector `age` and `participants`.

## Introduction to R objects

```
c(age, participants)
```

```
[1] "32"      "34"      "18"      "22"      "65"      "Paul"    "Maria"   "Andres"  
[9] "Robert"  "Alicia"
```

- What's the problem with this result?

# This breaks an R law!

We joined a numeric vector and a character vector.

Vectors can *ONLY* be of one class.

```
c(1, "one") # forces to character vector
```

```
[1] "1"  "one"
```

```
c(1, "1") # note that the first one is a numeric, while the second is a character
```

```
[1] "1" "1"
```



## Introduction to R objects

Now, which of these people has an age above 20?

```
age > 20
```

```
[1] TRUE TRUE FALSE TRUE TRUE
```

- That's a logical vector.

**Contrary to** character **and** numeric **vectors**, logical **vectors** can only have three values:

- TRUE
- FALSE
- NA (which stands for “Not available”.)

## Introduction to R objects

logicals can be created manually or with a logical statement.

```
c(TRUE, FALSE, TRUE, TRUE)
```

```
[1] TRUE FALSE TRUE TRUE
```

```
age < 60
```

```
[1] TRUE TRUE TRUE TRUE FALSE
```

The above expression tests for the logical statement.

For example,

```
  32    34    18    22    65  
TRUE TRUE TRUE TRUE FALSE
```

## Introduction to R objects

You can also write `T` or `F` as short abbreviations of `TRUE` and `FALSE`.

```
c(T, T, F, T) == c(TRUE, TRUE, FALSE, TRUE)
```

```
[1] TRUE TRUE TRUE TRUE
```

Which is comparing:

TRUE	TRUE	FALSE	TRUE
"T"	"T"	"F"	"T"

But behind the scenes, `TRUE` and `T` are just a `1` and `F` and `FALSE` are just a `0`.

What is the result of this?

```
T + 5  
TRUE - 5  
FALSE + TRUE  
T + T - FALSE
```

## Introduction to R objects

Now that you know that.. what would be the class of the following vectors?

```
c(5, TRUE)
c(5, "FALSE")
c(FALSE, TRUE)
c(1, FALSE)
```

- numeric: TRUE is coerced to 1
- character: "FALSE" is a string, can't be turned to a number
- logical: both elements are logical!
- numeric: FALSE is coerced to 0

## Introduction to R objects

We used some logical evaluations on those previous slides (*Is age greater than 20?*). But what is the list of operators available?

Logical evaluations (generally) work like arithmetic ones, that is, the evaluation is executed element-wise:

```
age > 0 # Is each value greater than zero?  
age == 22 # Are corresponding values equivalent?  
age != 22 # Are corresponding values not equivalent?
```

However, other logical evaluations take the whole object as the element to evaluate:

```
"z" %in% letters # Is "z" present in the list letters?
```

```
[1] TRUE
```

## Introduction to R objects

### What do we know so far?

- Numeric vectors
- Character vectors
- Logical vectors
- How to assign a name to these vectors
- Vectors can contain only *one* class of data

### What's missing?

#### **Factors**

## Introduction to R objects

Factors are R's way of storing **categorical variables**.

Categories such as:

- 'Male' and 'Female' or 'Married' and 'Divorced'
- 'Good', 'Middle' and 'High'

```
gender <- c("Male", "Female", "Male", "Male", "Female")  
  
# Can be turned into  
gender <- factor(gender)
```

## Introduction to R objects



## Introduction to R objects

Factors are useful for some specific operations like:

- Changing order of levels for terms in modelling
- Changing order of axis labels in plots
- Among other things..

In many cases you can use characters to do what you would want with factors!

## Introduction to R objects

Now, have you noticed that we've been assigning names to things?

```
age
```

```
[1] 32 34 18 22 65
```

The name `age` holds all these elements inside. How do we know where all the variables we've created are?

Let's ask R what objects can be listed from our workspace or environment.

```
ls()
```

```
[1] "age"      "elm"      "elm_factor" "gender"    "lgl"  
[6] "participants"
```

## Introduction to R objects

So far, we have a bunch of variables scattered around our workspace. This is usually no the way to go!

**We want to group similar things in the same place.**

```
our_df <- data.frame(name = participants, age = age, gender = gender, age_60 = lgl)  
our_df
```

	name	age	gender	age_60
1	Paul	32	Male	TRUE
2	Maria	34	Female	TRUE
3	Andres	18	Male	TRUE
4	Robert	22	Male	TRUE
5	Alicia	65	Female	FALSE

A data frame is usually the primary structure of analysis in R

## Introduction to R objects

It's important that you understand the thing that defines a data frame.

- A data frame has *rows* and *columns*, more technically called *dimensions*.
- Data frames have two dimensions.

```
dim(our_df)
```

```
[1] 5 4
```

```
nrow(our_df)
```

```
[1] 5
```

```
ncol(our_df)
```

```
[1] 4
```

## Introduction to R objects

Data frames are very distinctive because they can hold any type of vector.

Matrices cannot!

```
our_matrix <- matrix(1:20, ncol = 4, nrow = 5)
our_matrix
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	6	11	16
[2,]	2	7	12	17
[3,]	3	8	13	18
[4,]	4	9	14	19
[5,]	5	10	15	20

- Matrices are very similar to data frames.
- They have same number of dimensions.
- You can choose rows/columns in similar ways.

Introduction to R objects

Finally, we're missing the secret ingredient that differentiates both matrices and data frames.

## Lists

## Introduction to R objects

Think of lists as a bag that can store anything.

```
our_list <- list(names = participants, gender = gender, age = age)
our_list
```

```
$names
[1] "Paul"    "Maria"   "Andres"  "Robert"  "Alicia"

$gender
[1] Male    Female Male    Male    Female
Levels: Female Male

$age
[1] 32 34 18 22 65
```

This is a bag that has 3 objects.

- A character
- A factor
- A numeric

## Introduction to R objects

Think outside the box... when I say anything, I mean ANYTHING!


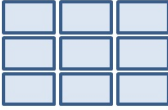


```
complex_list <- list(df = our_df[1:3, ], matrix = our_df[1:3, ], avg_age = mean(age))  
complex_list
```

```
$df  
  name age gender age_60  
1  Paul  32   Male   TRUE  
2 Maria  34 Female   TRUE  
3 Andres 18   Male   TRUE  
  
$matrix  
  name age gender age_60  
1  Paul  32   Male   TRUE  
2 Maria  34 Female   TRUE  
3 Andres 18   Male   TRUE  
  
$avg_age  
[1] 34.2
```



## Introduction to R objects

To sum up, these are the 4 types of data structures available in R.

Vector	Matrix	Data Frame	List
			
<ul style="list-style-type: none"><li>- 1 column or row of data</li><li>- 1 type (numeric or text)</li></ul>	<ul style="list-style-type: none"><li>- multiple columns and/or rows of data</li><li>- 1 type (numeric or text)</li></ul>	<ul style="list-style-type: none"><li>- multiple columns and/or rows of data</li><li>- multiple types</li></ul>	<ul style="list-style-type: none"><li>- 1 column or row of data</li><li>- multiple types</li></ul>

## Introduction to R objects

Now I'm gonna rock your world...

A data frame is a list (because it can have any class) with a row and column dimensions.

```
data.frame(our_list)
```

	names	gender	age
1	Paul	Male	32
2	Maria	Female	34
3	Andres	Male	18
4	Robert	Male	22
5	Alicia	Female	65

To be continued....